

Some results for melspectrogram/ MFCC/ raw wave nets

1. raw wave (1d)

```
class Simple1dNet(nn.Module):
    def __init__(self, in_shape=(1,16000), num_classes=12 ):
        super(Simple1dNet, self).__init__()

        self.conv1a = ConvBn1d( 1, 8, kernel_size=3, stride=1)
        self.conv1b = ConvBn1d( 8, 8, kernel_size=3, stride=1)

        self.conv2a = ConvBn1d( 8, 16, kernel_size=3, stride=1)
        self.conv2b = ConvBn1d( 16, 16, kernel_size=3, stride=1)

        self.conv3a = ConvBn1d( 16, 32, kernel_size=3, stride=1)
        self.conv3b = ConvBn1d( 32, 32, kernel_size=3, stride=1)

        self.conv4a = ConvBn1d( 32, 64, kernel_size=3, stride=1)
        self.conv4b = ConvBn1d( 64, 64, kernel_size=3, stride=1)

        self.conv5a = ConvBn1d( 64,128, kernel_size=3, stride=1)
        self.conv5b = ConvBn1d(128,128, kernel_size=3, stride=1)

        self.conv6a = ConvBn1d(128,256, kernel_size=3, stride=1)
        self.conv6b = ConvBn1d(256,256, kernel_size=3, stride=1)
        self.conv7a = ConvBn1d(256,256, kernel_size=3, stride=1)
        self.conv7b = ConvBn1d(256,256, kernel_size=3, stride=1)

        self.conv8a = ConvBn1d(256,512, kernel_size=3, stride=1)
        self.conv8b = ConvBn1d(512,512, kernel_size=3, stride=1)
        self.conv9a = ConvBn1d(512,512, kernel_size=3, stride=1)
        self.conv9b = ConvBn1d(512,512, kernel_size=3, stride=1)
        #self.linear1 = nn.Linear(512*31,1024)

        self.conv10a = ConvBn1d( 512,1024, kernel_size=3, stride=1)
        self.conv10b = ConvBn1d(1024,1024, kernel_size=3, stride=1)

        self.linear1 = nn.Linear(1024,512)
        self.linear2 = nn.Linear(512,256)
        self.fc       = nn.Linear(256,num_classes)
```

```
def forward(self, x):

    #print(x.size())
    x = F.relu(self.conv1a(x),inplace=True)
    x = F.relu(self.conv1b(x),inplace=True)
    x = F.max_pool1d(x,kernel_size=2,stride=2)

    #print(x.size())
    x = F.relu(self.conv2a(x),inplace=True)
    x = F.relu(self.conv2b(x),inplace=True)
    x = F.max_pool1d(x,kernel_size=2,stride=2)

    #print(x.size())
    x = F.relu(self.conv3a(x),inplace=True)
    x = F.relu(self.conv3b(x),inplace=True)
    x = F.max_pool1d(x,kernel_size=2,stride=2)

    #print(x.size())
    x = F.dropout(x,p=0.10,training=self.training)
    x = F.relu(self.conv4a(x),inplace=True)
    x = F.relu(self.conv4b(x),inplace=True)
    x = F.max_pool1d(x,kernel_size=2,stride=2)



    #print(x.size())
    x = F.dropout(x,p=0.10,training=self.training)
    x = F.relu(self.conv5a(x),inplace=True)
    x = F.relu(self.conv5b(x),inplace=True)
    x = F.max_pool1d(x,kernel_size=2,stride=2)

    ■ ■ ■

    #print(x.size())
    x = F.adaptive_avg_pool1d(x,1)
    x = x.view(x.size(0), -1)
    x = F.dropout(x,p=0.50,training=self.training)
    x = F.relu(self.linear1(x),inplace=True)

    x = F.dropout(x,p=0.50,training=self.training)
    x = F.relu(self.linear2(x),inplace=True)
    x = self.fc(x)

    return x #logits
```



30

31 waves_per_epoch = 51088

32

33 rate iter_k epoch num_m | valid_loss/acc | train_loss/acc | batch_loss/acc | time

34

35 0.0000 0.0 k 0.00 0.0 | 2.4848 0.0998 | 0.0000 0.0000 | 0.0000 0.0000 | 0 hr 00 min

36 0.0100 0.5 k 2.51 0.1 | 1.2439 0.6186 | 1.2220 0.5924 | 1.1421 0.6523 | 0 hr 02 min

37 0.0100 1.0 k 5.01 0.3 | 0.6321 0.8082 | 0.5843 0.8195 | 0.6067 0.8086 | 0 hr 05 min

38 0.0100 1.5 k 7.52 0.4 | 0.3553 0.8871 | 0.4567 0.8590 | 0.4639 0.8555 | 0 hr 08 min

39 0.0100 2.0 k 10.02 0.5 | 0.2723 0.9183 | 0.4169 0.8672 | 0.3099 0.9023 | 0 hr 11 min

40 0.0100 2.5 k 12.53 0.6 | 0.2650 0.9133 | 0.3205 0.9049 | 0.3491 0.8984 | 0 hr 14 min

41 0.0100 3.0 k 15.03 0.8 | 0.2724 0.9171 | 0.2745 0.9158 | 0.3132 0.9062 | 0 hr 17 min

42 0.0100 3.5 k 17.54 0.9 | 0.2419 0.9258 | 0.2669 0.9203 | 0.2554 0.9219 | 0 hr 20 min

2. melspectrogram (2d)

```
class SeResNet3(nn.Module):
    def __init__(self, in_shape=(1,40,101), num_classes=12 ):
        super(SeResNet3, self).__init__()
        in_channels = in_shape[0]

        self.layer1a = ConvBn2d(in_channels, 16, kernel_size=(3, 3), stride=(1, 1))
        self.layer1b = ResBlock( 16, 16)

        self.layer2a = ConvBn2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
        self.layer2b = ResBlock(32, 32)
        self.layer2c = ResBlock(32, 32)

        self.layer3a = ConvBn2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
        self.layer3b = ResBlock(64, 64)
        self.layer3c = ResBlock(64, 64)

        self.layer4a = ConvBn2d( 64,128, kernel_size=(3, 3), stride=(1, 1))
        self.layer4b = ResBlock(128,128)
        self.layer4c = ResBlock(128,128)

        self.layer5a = ConvBn2d(128, 256, kernel_size=(3, 3), stride=(1, 1))
        self.layer5b = nn.Linear(256,256)

        self.fc = nn.Linear(256,num_classes)
```

```
class ConvBn2d(nn.Module):
    def __init__(self, in_channels, out_channels, kernel_size=3, padding=1, dilation=1, stride=1):
        super(ConvBn2d, self).__init__()
        self.conv = nn.Conv2d(in_channels, out_channels, kernel_size=kernel_size, padding=padding, dilation=dilation, stride=stride)
        self.bn = nn.BatchNorm2d(out_channels)
        if is_bn is False:
            self.bn = None

    def forward(self, x):
        x = self.conv(x)
        if self.bn is not None:
            x = self.bn(x)
        return x

class SeScale(nn.Module):
    def __init__(self, channel, reduction=16):
        super(SeScale, self).__init__()
        self.fc1 = nn.Conv2d(channel, reduction, kernel_size=1, padding=0)
        self.fc2 = nn.Conv2d(reduction, channel, kernel_size=1, padding=0)



    def forward(self, x):
        x = F.adaptive_avg_pool2d(x,1)
        x = self.fc1(x)
        x = F.relu(x, inplace=True)
        x = self.fc2(x)
        x = F.sigmoid(x)
        return x

class ResBlock(nn.Module):
    def __init__(self, in_planes, out_planes, reduction=16):
        super(ResBlock, self).__init__()
        assert(in_planes==out_planes)

        self.conv_bn1 = ConvBn2d(in_planes, out_planes, kernel_size=3, padding=1, stride=1)
        self.conv_bn2 = ConvBn2d(out_planes, out_planes, kernel_size=3, padding=1, stride=1)
        self.scale = SeScale(out_planes, reduction)

    def forward(self, x):
        z = F.relu(self.conv_bn1(x),inplace=True)
        z = self.conv_bn2(z)
        z = self.scale(z)*z + x
        z = F.relu(z,inplace=True)
        return z
```

```
def forward(self, x):  
    x = F.relu(self.layer1a(x), inplace=True)  
    x = self.layer1b(x)  
    x = F.max_pool2d(x, kernel_size=(2,2), stride=(2,2))  
  
    x = F.dropout(x, p=0.1, training=self.training)  
    x = F.relu(self.layer2a(x), inplace=True)  
    x = self.layer2b(x)  
    x = self.layer2c(x)  
    x = F.max_pool2d(x, kernel_size=(2,2), stride=(2,2))  
  
    x = F.dropout(x, p=0.2, training=self.training)  
    x = F.relu(self.layer3a(x), inplace=True)  
    x = self.layer3b(x)  
    x = self.layer3c(x)  
    x = F.max_pool2d(x, kernel_size=(2,2), stride=(2,2))  
  
    x = F.dropout(x, p=0.2, training=self.training)  
    x = F.relu(self.layer4a(x), inplace=True)  
    x = self.layer4b(x)  
    x = self.layer4c(x)  
  
    x = F.dropout(x, p=0.2, training=self.training)  
    x = F.relu(self.layer5a(x), inplace=True)  
    x = F.adaptive_avg_pool2d(x, 1)  
    x = x.view(x.size(0), -1)  
    x = F.relu(self.layer5b(x))  
  
    x = F.dropout(x, p=0.2, training=self.training)  
    x = self.fc(x)  
  
    return x #logits
```



LB@21k = ?
Valid = 0.1099 0.9716
Train = 0.0804 0.9777

LB@20k =
Valid = 0.1130 0.9688
Train = 0.0988 0.9719

waves_per_epoch = 57886

rate	iter_k	epoch	num_m	valid_loss/acc		train_loss/acc		batch_loss/acc		time
0.0000	0.0 k	0.00	0.0	2.4853	0.0808	0.0000	0.0000	0.0000	0.0000	0 hr 00 min
0.0100	0.5 k	2.21	0.1	1.1928	0.5332	1.3250	0.4996	1.0544	0.6133	0 hr 01 min
0.0100	1.0 k	4.42	0.3	0.4840	0.8665	0.3761	0.8875	0.4057	0.8711	0 hr 02 min
0.0100	1.5 k	6.63	0.4	0.1921	0.9370	0.2454	0.9273	0.3459	0.8945	0 hr 03 min
0.0100	2.0 k	8.84	0.5	0.1868	0.9379	0.2417	0.9295	0.2735	0.9102	0 hr 04 min
0.0100	2.5 k	11.06	0.6	0.1574	0.9461	0.2103	0.9418	0.2307	0.9375	0 hr 06 min
0.0100	3.0 k	13.27	0.8	0.1575	0.9523	0.1911	0.9449	0.2082	0.9336	0 hr 07 min
0.0100	3.5 k	15.48	0.9	0.1688	0.9470	0.1840	0.9453	0.1646	0.9492	0 hr 08 min
0.0100	4.0 k	17.69	1.0	0.1489	0.9526	0.1603	0.9543	0.1255	0.9648	0 hr 09 min
0.0100	4.5 k	19.90	1.2	0.1415	0.9620	0.1627	0.9514	0.1444	0.9531	0 hr 11 min
0.0100	5.0 k	22.11	1.3	0.1612	0.9479	0.1519	0.9570	0.1685	0.9531	0 hr 12 min
0.0100	5.5 k	24.32	1.4	0.1058	0.9669	0.1424	0.9607	0.1294	0.9688	0 hr 13 min
0.0100	6.0 k	26.53	1.5	0.1056	0.9666	0.1440	0.9537	0.1208	0.9688	0 hr 14 min
0.0100	6.5 k	28.75	1.7	0.1087	0.9673	0.1328	0.9645	0.1358	0.9570	0 hr 15 min
0.0100	7.0 k	30.96	1.8	0.1339	0.9607	0.1404	0.9584	0.0958	0.9766	0 hr 17 min
0.0100	7.5 k	33.17	1.9	0.1109	0.9648	0.1331	0.9629	0.1499	0.9414	0 hr 18 min
0.0100	8.0 k	35.38	2.0	0.1076	0.9691	0.1330	0.9621	0.0753	0.9844	0 hr 19 min
0.0100	8.5 k	37.59	2.2	0.1259	0.9635	0.1283	0.9604	0.1998	0.9414	0 hr 20 min

3. mfcc(2d)

```
class VggNet1(nn.Module):
    def __init__(self, in_shape=(1,40,101), num_classes=12 ):
        super(VggNet1, self).__init__()

        self.conv1a = ConvBn2d( 1, 8, kernel_size=(3, 3), stride=(1, 1))
        self.conv1b = ConvBn2d( 8, 8, kernel_size=(3, 3), stride=(1, 1))
        self.conv2a = ConvBn2d( 8, 16, kernel_size=(3, 3), stride=(1, 1))
        self.conv2b = ConvBn2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
        self.conv3a = ConvBn2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
        self.conv3b = ConvBn2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        self.linear1 = nn.Linear(32*12*5,512)
        self.linear2 = nn.Linear(512,256)
        self.fc      = nn.Linear(256,num_classes)

    def forward(self, x):

        x = F.relu(self.conv1a(x),inplace=True)
        x = F.relu(self.conv1b(x),inplace=True)
        x = F.max_pool2d(x,kernel_size=(2,2),stride=(2,2))



        x = F.dropout(x,p=0.2,training=self.training)
        x = F.relu(self.conv2a(x),inplace=True)
        x = F.relu(self.conv2b(x),inplace=True)
        x = F.max_pool2d(x,kernel_size=(2,2),stride=(2,2))

        x = F.dropout(x,p=0.2,training=self.training)
        x = F.relu(self.conv3a(x),inplace=True)
        x = F.relu(self.conv3b(x),inplace=True)
        x = F.max_pool2d(x,kernel_size=(2,2),stride=(2,2))

        x = F.dropout(x,p=0.2,training=self.training)
        x = x.view(x.size(0), -1)
        x = F.relu(self.linear1(x),inplace=True)

        x = F.dropout(x,p=0.2,training=self.training)
        x = F.relu(self.linear2(x),inplace=True)
        x = self.fc(x)

        return x #logits
```



LB@83k =
Valid = 0.1821 0.9532
Train = 0.1803 0.9488

waves_per_epoch = 51088

rate	iter_k	epoch	num_m	valid_loss/acc	train_loss/acc	batch_loss/acc	time
0.0000	0.0 k	0.00	0.0	2.4852 0.0817	0.0000 0.0000	0.0000 0.0000	0 hr 00 min
0.0100	0.5 k	1.25	0.1	0.9850 0.6660	1.2023 0.5926	1.1932 0.5938	0 hr 01 min
0.0100	1.0 k	2.51	0.1	0.6628 0.7852	0.8045 0.7359	0.7815 0.7500	0 hr 03 min
0.0100	1.5 k	3.76	0.2	0.4714 0.8416	0.6400 0.7953	0.7717 0.7578	0 hr 04 min
0.0100	2.0 k	5.01	0.3	0.5584 0.8138	0.6133 0.7992	0.6682 0.7500	0 hr 06 min
0.0100	2.5 k	6.26	0.3	0.4070 0.8622	0.5415 0.8172	0.5128 0.8125	0 hr 08 min
0.0100	3.0 k	7.52	0.4	0.3844 0.8700	0.5453 0.8254	0.6335 0.7891	0 hr 09 min
0.0100	3.5 k	8.77	0.4	0.3471 0.8828	0.4641 0.8527	0.5630 0.8281	0 hr 11 min
0.0100	4.0 k	10.02	0.5	0.3392 0.8934	0.4467 0.8578	0.4552 0.8438	0 hr 13 min
0.0100	4.5 k	11.27	0.6	0.2808 0.9043	0.4337 0.8629	0.5054 0.8203	0 hr 14 min
0.0100	5.0 k	12.53	0.6	0.2796 0.9068	0.4043 0.8723	0.4035 0.8750	0 hr 16 min
0.0100	5.5 k	13.78	0.7	0.2816 0.9024	0.4073 0.8703	0.5822 0.8125	0 hr 17 min
0.0100	6.0 k	15.03	0.8	0.3261 0.8868	0.4005 0.8691	0.3281 0.8594	0 hr 19 min
0.0100	6.5 k	16.29	0.8	0.2524 0.9124	0.4138 0.8691	0.3202 0.9141	0 hr 21 min
0.0100	7.0 k	17.54	0.9	0.2599 0.9180	0.3840 0.8762	0.3712 0.8672	0 hr 22 min
0.0100	7.5 k	18.79	1.0	0.2400 0.9205	0.4027 0.8742	0.5240 0.8281	0 hr 24 min
0.0100	8.0 k	20.04	1.0	0.2436 0.9199	0.3603 0.8888	0.4499 0.8438	0 hr 26 min
0.0100	8.5 k	21.30	1.1	0.2336 0.9255	0.3585 0.8887	0.4363 0.8750	0 hr 27 min
0.0100	9.0 k	22.55	1.2	0.2169 0.9317	0.3959 0.8691	0.2485 0.9141	0 hr 29 min